

JAYESH YADAV

Bengaluru, Karnataka

701-494-3090 | 007jayeshyadav@gmail.com | linkedin.com/in/007jayesh | medium.com/@007jayeshyadav

Education

Indian Institute of Technology, Madras

CGPA: 8.1

Diploma in Data Science

University College of Engineering and Technology, RTU

CGPA: 7.4

Bachelor of Science in Computer Science

Experience

Quantifai.ai – AI-Powered Reconciliation & Finance Automation

Feb 2025 – Present

Bengaluru, India

Data & AI Engineer

- **VC Fund Reconciliation Engine:** Built end-to-end document processing pipeline for capital call PDFs and bank statements (PDF/Excel) with 100% accuracy requirement. Designed Docling-first + LLM fallback architecture, reducing API costs by 70% while processing 500+ documents/month. All extraction enforced through Pydantic validation models.
- **NER-Based Transaction Matching:** Built investor-name matching algorithm achieving 0% false positives across 10,000+ transactions with human-in-the-loop. Developed adaptive learning system that records user corrections into a knowledge base and auto-matches after 3 confirmations – reducing manual intervention by 40% MoM.
- **Browser-to-API Platform (Inspired by Parse.bot):** Built proprietary engine that converts any web portal into a programmable API, auto-generating Python SDKs from browser interactions. Added email ingestion agent for vendors distributing settlement data via email. Powering 30+ vendor integrations, processing 1M+ transactions/day, saving 150+ engineering hours/month. Reusable connectors cut onboarding from 2 weeks to 2 days.
- **Data Fabric – Self-Healing Pipelines (In Progress):** Building zero-code onboarding where users upload any document and system auto-generates parsers, tests in sandbox, and deploys to backend. Built AI Agent that auto-detects schema drift, self-heals pipelines, and sends root-cause alerts for approval. Includes natural-language transformation rules engine eliminating developer dependency for business rule changes.
- **Reconciliation Dashboard & Reporting:** Built interactive exception management dashboard with store-wise drill-downs, anomaly detection, and one-click audit-ready report generation – reducing month-end close from 10+ days to 3 days for multi-outlet retail clients.
- **Platform Impact:** End-to-end ownership of all AI/data infrastructure – from document ingestion to reconciliation output. Platform serves F&B chains, VC funds, and retail businesses with fully automated reconciliation workflows.

Log9 Materials – EV Battery Technology

Mar 2023 – Jan 2025

Bengaluru, India

Data & AI Engineer

- **CellView – Automated Battery Analytics Platform:** Built full-stack analytics platform comparable to Voltaiq (\$200K+/yr SaaS), saving \$200K+ annually. Engineered automated data pipelines from battery cyclers (4-6 hrs/day → real-time). Built LLM-powered no-code dashboard for natural language analytics. Adopted by 3 departments, 25+ DAU, 500K+ data points/day from 50+ cyclers.
- **SOH Prediction Model – EV Insurance:** Built first-of-its-kind battery health prediction model (98% R²) enabling insurance companies to replace engine-based premium calculations. Deployed across 1,200+ vehicles, contributing to \$150K ARR partnership with insurance providers.
- **Battery Pack Viability Prediction:** Meta-model (Random Forest + BiLSTM) predicting pack pass/fail before assembly with R² of 0.93, reducing testing time from 8 hrs to 15 min/batch, cutting rework by 35% and saving \$80K/year.
- **AutoML & Automated Reporting:** Integrated AutoML module for predictive modeling of battery degradation patterns, cutting QC cycle time by 60%. Automated daily stakeholder reports with department-specific views, reducing decision turnaround from 3 days to same-day.
- **Org Impact:** Operated as a one-person data & AI team across R&D, Quality, and Production. Combined savings across all projects exceeded \$350K/year in tooling costs, manual labor, and material waste reduction.

KPMG

Nov 2022 – Mar 2023

Data Analytics Consulting Intern

- Rectified 15,000+ data inconsistencies across 3 client datasets (integrity: 72% → 96%). Built automated pipelines in Python/SQL reducing reporting turnaround from 5 days to 1 day, delivering actionable insights for strategic decision-making.

Entrepreneurial Project

Neuracer – Procurement Automation SaaS

Dec 2025 – Present

Bootstrapped startup | Full-stack AI-assisted development (Cursor, Claude Code)

- Building an end-to-end procurement automation platform for SMB manufacturers featuring three-way matching (Purchase Orders ↔ Invoices ↔ Goods Receipts), LLM-powered document extraction with GST compliance, and ERP integrations via API/XML-RPC.
- Architected and shipped the entire product using AI-assisted development, eliminating the need for a dedicated frontend/backend developer. Multi-tenant PostgreSQL design with JSONB custom fields for per-client configurability.

Technical Skills

Languages & Frameworks: Python, SQL, FastAPI, Streamlit, Pydantic, React (Vibe Coding)

AI/ML: LLM Pipelines (Gemini, Claude), NER, CatBoost, BiLSTM, AutoML, RAG, AI Agents

Data & Automation: Docling, Browser-to-API, Reverse Engineering, ETL, Self-Healing Pipelines, Selenium

Tools: Git, Docker, PostgreSQL, Linux, Cursor, Claude Code, Pandas, NumPy, Plotly, Odoo ERP